

Marketing & Launch Plan

— built to be remembered, not just launched

Owner: Sanjoy Ghosh · **Last updated:** 2026-07-11 · Revalidates the 2026-07-10 v1 plan against one explicit lens: durable memorability, zero cost, collaboration/feedback velocity, and personal brand compounding. Paste this whole page — or the source at [MARKETING.md](#) — at the start of any future session to resume with full context.

WHAT CHANGED IN V2

① A flagship asset

v1 was all "post about the tool." v2 adds something worth sharing on its own — a recurring public model benchmark.

② A collaboration home

Discussions, good-first-issue labels, contributor recognition — v1 never named where engagement actually lands.

③ A brand track, not just product

A second content stream that builds your name specifically, independent of any single launch spike.

④ Two new channels

Reactive newsjacking and short-form video — zero cost, both entirely missing from v1.

What changed since v1 — the honest revalidation

v1 was a solid, correctly-scoped channel checklist: accurate HN/PH mechanics, a sensible audience order, free-only discipline throughout. It would generate a respectable launch-day spike. It was not, on its own, built to be *remembered* — there was no single asset designed to be inherently shareable independent of "here's a new tool," no explicit collaboration home, and no track for your own name to compound past this one launch. Those are the gaps this revision closes.

Everything else from v1 — the four narrative angles, the audience priority order, Tier 1–3 channels, the 90-day skeleton, the capability split, the metrics — held up and is carried forward with small factual corrections, marked inline below.

The single biggest miss in v1

Every channel entry answers "where do we post about GAUNTLEX," but none of them answer "what does GAUNTLEX *produce* that people want to share on its own." Tools get a launch-day spike. Ongoing, original research gets cited, linked, and re-shared for months — that's the difference between a good Show HN and something people still reference next year.

The flagship asset — give people something to share, not just read

2.1 The AI Model Resilience Index — **primary recommendation**

The idea: GAUNTLEX already scores how well generated code survives adversarial testing. Point that at the question every developer in this space is quietly wondering right now: *which AI model writes the most secure code by default?* Run the same handful of intentionally risk-prone specs — an auth endpoint, a file-upload handler, a raw SQL query builder, the classic OWASP-magnet patterns — through several current models as the Builder, score each resulting implementation with GAUNTLEX's own Arbiter, and publish the ranked result as a public leaderboard with a short written analysis.

Why this is the right flagship, specifically now

It's original research, not a product pitch. "We built a tool" is forgettable. "GPT-5-class models missed 3x more injection attacks than Claude in our test" is a stat people screenshot and argue about. That argument *is* the buzz.

Comparison is the most reliably shareable content format there is. Rankings and "X vs Y" posts outperform feature announcements on every channel in this plan — HN, Twitter, Reddit, LinkedIn — by a wide margin, consistently.

It's recurring, not one-off. Re-run it monthly as models update ("Resilience Index — August 2026 edition") and you have a standing reason to post again, get cited again, and pull people back — the single highest-leverage fix for "remembered for a long time" in this entire plan.

It costs nothing. OpenRouter's free tier alone covers several current open models; comparing what's reachable for free is itself a defensible, honest framing.

It directly demonstrates the product — every visitor reading the leaderboard is looking at real GAUNTLEX output, no gap between the marketing claim and the tool.

It's ethically clean. You're scoring *models* — products from well-resourced companies fully expecting to be benchmarked — not a stranger's weekend project, which would read as punching down.

 CLAUDE DOES

Design 2–3 shared adversarial specs. Run each through a defensible list of free/accessible OpenRouter models as the Builder. Score every implementation with the existing Arbiter. Generate the static leaderboard via the `gauntlet leaderboard` command that already exists. Write the analysis post. Match the existing dashboard theme.

 YOU DO

Approve the exact model list before it's public — comparing named commercial products is a claim that should get your eyes on it. Decide the cadence (suggest monthly). Post the launch of it, same as any other content.

Sequencing: build it in Weeks 1–2, *before* Show HN and Product Hunt — so those launch posts link to a live, populated leaderboard instead of just a README. This is the single highest-leverage sequencing change in this revision.

2.2 "GAUNTLEX red-teams itself" – secondary, near-zero effort

A one-line but disproportionately effective trust signal: run GAUNTLEX against its own specification and link the resulting Resilience Report — SHA-256 integrity hash and all — directly in the README and every launch post. Security tools that market themselves without receipts read as marketing; one that hands you a tamper-evident report on itself reads as credible.

 CLAUDE DOES

Run it, publish the report link, add a "we tested ourselves — verify it yourself with `gauntlet verify`" line to the README and launch copy.

 YOU DO

Sanity-check the result before it's public. Nothing else.

Community & collaboration loop

You asked for people to collaborate, give feedback, engage — not just install. v1 never actually designated *where* that happens. Fixing that:

1 Enable GitHub Discussions

Make it the one explicit home for feedback, questions, and show-and-tell — link it from the README header, every social post's CTA, and the HN/PH copy. Right now there is no owned space for conversation to land; issues aren't it — treating them as the general forum drives away casual feedback.

2 Seed 3–5 real "good first issue" labels

Approachable, genuinely useful work — a missing language profile, a policy domain gap, a docs improvement. A repo with zero labeled entry points silently tells contributors "don't bother."

3 A visible Contributors table

Even a short one, updated as people show up. Recognition is the cheapest retention mechanic that exists and costs nothing to maintain.

4 Every piece of content ends with one specific, low-friction ask

Not "check it out" — which asks for nothing and gets nothing — but something concrete: "reply with the worst security bug your AI coding assistant ever shipped you." Specific asks get specific responses; vague CTAs get silence.

5 The "Powered by GAUNTLEX" badge

Carried over from v1 — every adopting repo becomes a small, permanent ad. Worth prioritizing earlier than v1 suggested: it's a collaboration mechanic, not just a growth one — visible proof other people trust this.

CLAUDE DOES

Enable Discussions via `gh` CLI with your approval. Draft good-first-issue candidates and label them. Build the Contributors table. Bake the specific-CTA discipline into every drafted post.

YOU DO

Approve enabling Discussions. Actually respond in there once people show up — this is the part that can't be delegated.

Personal brand track — a second, parallel content stream

v1's entire content calendar is product-news-shaped: launch posts, feature posts, "we shipped this" posts. That builds awareness of GAUNTLEX. It does comparatively little for *you*, specifically, being the person people think of when this topic comes up — and that recognition is what compounds past this launch into whatever you build next.

Run two tracks, not one

TRACK A — PRODUCT

v1's existing plan

Launch posts, feature announcements, the Resilience Index.

Cadence: tied to ship events.

→ builds GAUNTLEX awareness

TRACK B — FIELD NOTES

New

Short, opinionated, single-observation posts on AI-generated code security — real patterns you're seeing. The tool shows up as evidence for a claim, not the headline.

Cadence: ~weekly, decoupled from any release.

→ builds *your* authority, independent of GAUNTLEX awareness

Field Notes is the more durable of the two. A launch post's relevance has a half-life of days; a genuinely sharp observation about a real pattern in AI-generated code gets referenced any time the topic comes up, for as long as the pattern stays true. It's also lower-effort once running — it doesn't require a ship event to justify a post, so it fills the gaps between the launch-driven calendar.

A rallying line

Worth having one short, ownable phrase people repeat independent of the tool name — the way "shift left" outlived its origin. Three directions, your call on tone — I can't invent this authentically for you, same as the origin story:

"Code ships in seconds. Attacks should too."

"If the AI wrote it in 30 seconds, the attack surface should be tested in 30 seconds."

"Adversarial by default."

 CLAUDE DOES

Draft the Field Notes series — mining real patterns from the AVF golden fixtures, the CWE taxonomy, and general AI-code-security research. Propose the rallying-line options above.

 YOU DO

Pick or rewrite the rallying line. Add your actual voice/opinion to each Field Notes draft — this track specifically fails if it reads as AI-generated. Post everything yourself.

§5 · New

Reactive coverage — newsjacking, done carefully

When a real AI-generated-code security incident or CVE breaks publicly, being one of the first credible voices connecting it to your exact thesis is disproportionately effective compared to any scheduled post — it rides existing attention instead of trying to create its own.

 CLAUDE DOES

On request, monitor for relevant incidents via web search and flag candidates with a one-paragraph brief on the angle.

 YOU DO

The actual reactive post has to be fast and has to be yours — by the time a draft goes through review-and-approve, the moment's usually gone. Keep this on-demand, not scheduled.

Short-form video

You already have a 2:43 demo video built from real screenshots. A 30–45 second vertical cut of the single most visual moment — Builder and Breaker racing concurrently, or the live ARS gate blocking a merge — is a genuinely different discovery surface: YouTube Shorts, TikTok, Instagram Reels, and LinkedIn video all have algorithmic reach that doesn't depend on an existing following. Free, and it reuses an asset you've already paid the production cost for.

CLAUDE DOES

Identify the strongest 30–45s segment from the existing footage, cut and caption it (burned-in — most short-form video is watched muted), format vertical (9:16).

YOU DO

Upload to your own accounts — same authenticated-human-presence constraint as everything else.

Positioning – carried over from v1, unchanged

Core thesis: "AI coding tools generate code faster than any human ever could. Nobody generates the security testing at the same speed — it's still bolted on after, by a human, if at all. GAUNTLEX closes that gap: it generates the attack at the exact same instant it generates the code."

ANGLE	BEST FOR	ONE-LINER
The speed gap	HN, Twitter, general dev	"Your AI pair programmer ships code in 30 seconds. Your security review takes 3 days. GAUNTLEX makes them the same speed."
TDD for security	Engineering managers, LinkedIn	"You wouldn't ship code without tests. Why ship AI-generated code without adversarial tests?"
Compliance-as-code	FinTech/HealthTech, LinkedIn, enterprise	"HIPAA and FINRA audits shouldn't be a fire drill every quarter."
Novel architecture	HN, r/MachineLearning, research	"Builder and Breaker run concurrently via <code>asyncio.gather()</code> — the attack surface is reasoned from the spec alone."
The Resilience Index	Everywhere — this one has a hard number attached	"We ran the same spec through N AI models and scored what came back. Here's which ones actually write secure code by default."

Your personal story still matters more than the product spec — unchanged, and still the single most important missing input (§14).

Target audiences — carried over, unchanged

1. **AI-native developers** — living in Claude Code / Cursor / Copilot / Codex, worried (or should be) about what their AI just shipped
2. **AppSec / DevSecOps engineers** — the people whose job GAUNTLEX makes easier, not harder
3. **Engineering leadership** at fast-moving startups shipping AI-generated code without a security net
4. **Compliance-heavy orgs** — FinTech, HealthTech, anyone who has said "SOC 2" or "HIPAA" in a board meeting
5. **MCP ecosystem builders** — smaller but extremely high-intent, actively browsing MCP servers right now

Asset inventory – corrected as of 2026-07-11

ASSET	STATUS	NOTES
Public GitHub repo, MIT license	have	github.com/sanjoy1234/gauntlex
PyPI package (gauntlex-ai)	have	Live at v1.0.1 — published 2026-07-10, verified installable
Demo video (2:43)	have	Source for the new short-form cuts, §6
README, badges, quickstart	have	Test-count badge now current (612 passing)
Live dashboard + leaderboard	have	The exact engine the Resilience Index (§2.1) reuses
AI Model Resilience Index	to build	New — launch centerpiece, §2.1
GitHub Discussions	to enable	Not currently on — §3, needs your approval
Show HN / Product Hunt copy	to draft	§10.5, §10.6
Twitter/LinkedIn content	to draft	Now split: Product + Field Notes, §4
Short-form video cuts	to produce	New, §6
Rallying line	need your pick	Three options drafted, §4
Your origin story	need from you	Still the single most important missing input
Social accounts (X, LinkedIn, dev.to, PH, Reddit, YouTube/TikTok)	confirm	Tell me which exist already

Channel playbook

v1's tier structure, tactics, and platform-specific rules were reviewed against current information and are unchanged — reproduced here so this page is the single reference going forward.

Tier 1 — do these first

Cost free **Reach** highly targeted **Effort** low

DIRECTORY	NOTES
Official MCP Registry	registry.modelcontextprotocol.io — reverse-DNS naming, CLI publish
Glama	glama.ai/mcp — auto-indexes public MCP servers, verify/claim your listing
Smithery	<code>smithery mcp publish</code>
PulseMCP	18,000+ servers, updated daily
mcp.so	Large community directory
awesome-mcp-servers (wong2 / punkpeye)	PR to add an entry, each
awesome-devops-mcp-servers	DevSecOps angle fits directly
Awesome-AI-For-Security / Awesome-LLMSecOps	Security-focused awesome lists

 **CLAUDE DOES**

Draft registry metadata, run CLI publish where I have shell access, open PRs matching each list's format.

 **YOU DO**

Approve each PR before it opens, claim the Glama auto-listing, complete any email verification.

tl;dr sec & TLDR newsletters

tl;dr sec: 90,000+ AppSec engineers, curated by Clint Gibler (Head of Cyber at OpenAI) — arguably your single best-fit audience anywhere. TLDR AI/DevOps/Infosec: 300–400K readers per edition.

 CLAUDE DOES

Draft tailored, no-hype submission pitches as Gmail drafts — one per edition, different angle each.

 YOU DO

Locate the current submission form/address, review, send.

Show HN

Title format: `Show HN: GAUNTLEX – Generates code and adversarial security tests at the same instant`. No superlatives. Post Tuesday–Thursday, 9 AM–12 PM ET. Never ask for upvotes. **HN bans AI-generated/AI-edited comments** — every reply must be written by you, live. Respond in the first 60 minutes — this drives ranking.

 CLAUDE DOES

Draft post title + body, link to the Resilience Index and demo video, prep anticipated objections.

 YOU DO

Post from your own account with real karma, write every comment reply yourself for at least 2 hours.

Product Hunt

Free to submit and be featured. Submit 12:00–1:00 AM PST, Tuesday or Wednesday. Self-hunting is fine. Tagline ≤60 characters, no emojis in the name field. Keep 1–2 weeks apart from Show HN so each gets its own wave.

 CLAUDE DOES

Draft tagline, description, maker comment, design 3–5 gallery images including Resilience Index screenshots.

 YOU DO

Confirm your maker profile isn't brand-new, submit at the right time, reply all day.

Tier 2 — ongoing

CHANNEL	NOTES
Reddit	r/programming, r/netsec, r/cybersecurity, r/devops, r/MachineLearning, r/ClaudeAI, r/cursor, r/ExperiencedDevs — tailored copy per subreddit, verify self-promo rules before posting
Twitter/X	Content calendar across both tracks (§4) — architecture deep-dive, origin story, launch thread, weekly Field Notes
LinkedIn	Compliance-as-code for Product track; Field Notes runs here too
dev.to / Hashnode	4 posts: origin+architecture, ARS deep-dive, compliance-in-CI, Resilience Index writeup
GitHub Discoverability	Topics (ai-security, mcp-server, devsecops, llm-security, adversarial-testing, sast), enable Discussions

Tier 3 — needs your voice/presence

CHANNEL	NOTES
Podcasts	Latent Space (best fit, confirmed to accept pitches), Application Security Podcast, The Changelog, Software Engineering Daily
Conference / Meetup CFPs	Local BSides, OWASP chapters, Python/AI meetups
"Powered by GAUNTLEX" badge	Reframed in §3 as a collaboration mechanic, not just growth

What I can build for you directly

Unchanged capability list from v1, still accurate:

- Draft every piece of written copy in this plan
- Create actual Gmail drafts in your inbox, ready to review and send
- Design social graphics and OG cards (HTML → headless-Chrome screenshot)
- Run the AI Model Resilience Index end-to-end — spec design, multi-model generation, scoring, publishing
- Open GitHub PRs to awesome-lists and registries via `gh` CLI, with your approval on each
- Research current submission processes so we're never guessing at a stale process
- Track this document and its Log as things ship

What I categorically cannot do

Post to your personal accounts, DM anyone as you, or take any action requiring your authenticated human presence on a third-party platform.

Revised 90-day sequencing

Weeks 1–2 — Foundation, no public launch risk yet

MCP registries + awesome-list PRs · GitHub Topics + enable Discussions · **build v1 of the AI Model Resilience Index** · finalize origin story + rallying line · draft all Tier 1 copy

Weeks 3–4 — Content warm-up, both tracks running

First dev.to post · Field Notes track starts on LinkedIn + Twitter, no launch framing yet · submit to tl;dr sec + TLDR · first short-form video clip

Week 5 — Show HN

Tuesday/Wednesday morning ET, links directly to the live Resilience Index · full-day availability

Weeks 6–7 — Product Hunt

Separate week from HN · coordinated LinkedIn/Twitter push same day · gallery includes Resilience Index screenshots

Weeks 8–12 — Sustained

Reddit posts spaced out · podcast pitches sent · second/third dev.to posts · **first Resilience Index refresh** — the moment that tests whether the recurring-asset thesis works · start tracking which channel actually drove stars/installs

Metrics to track

- GitHub stars/forks — leading indicator, easy to game, don't over-index
- PyPI download count — better signal than stars
- Referrer traffic on the repo (GitHub Insights → Traffic tab, free)
- HN/PH position and comment engagement on launch days
- Newsletter/podcast placements landed — binary, did it run or not
- **New** — Resilience Index citations/backlinks: the specific signal that tells you whether the flagship-asset thesis in §2 is working

What I need from you — the complete, minimal list

Everything else in this plan I can execute or draft without waiting on you. These six are the actual blockers.

1 Your one-sentence origin story

What made you build this. Still the single most important input — I can't invent it authentically.

2 Which accounts you already have

Twitter/X, LinkedIn, dev.to, Product Hunt maker profile, Reddit, YouTube/TikTok.

3 Sign-off on the Resilience Index model list

I'll propose a default set of free/accessible models — a public claim comparing named commercial products should get your eyes on it before it ships.

4 Approval to enable GitHub Discussions

One repo setting, done via `gh` CLI once you say go.

5 Your pick (or rewrite) of a rallying line

Three options drafted in §4, or tell me it's not you and we skip it.

6 Your voice on the Field Notes drafts

I'll write the first-pass observations, but this track fails if it doesn't sound like you.

Tell me these and I start building — the Resilience Index and the Tier 1 registry/awesome-list submissions can begin immediately, before any of the others are resolved.

Change log

2026-07-10 — v1 plan created.

2026-07-11 — v2: revalidated against memorability/collaboration/personal-brand lens. Added the AI Model Resilience Index as flagship asset (§2), a Community & Collaboration Loop (§3), a parallel Personal Brand / Field Notes track (§4), reactive-coverage practice (§5), and short-form video as a new channel (§6). Corrected asset inventory to current shipped state (PyPI v1.0.1 live, 612 tests, Discussions not yet enabled). Tightened "what I need from you" to six concrete blockers.
